

Socially-Weighted Distributed Graph Optimization (D²RO) for Autonomous Multi-Agent Service Fleets in Crowded Environments

Polla Fattah and Sanar Fawzi

Abstract—The continuous routing of autonomous service fleets—such as retail shopping carts (*Int-Cart*), clinical hospital pushchairs, and airport luggage trolleys—in crowded, human-shared public environments poses fundamental multi-agent coordination challenges. Classical shortest-path and reactive collision-avoidance methods treat pedestrians as obstacles to be cleared rather than as people whose personal space carries a cost, and they possess no mechanism for resolving symmetric conflicts in single-file corridors. This paper proposes the Distributed Dynamic Route Optimization (D²RO) framework powered by Socially-Weighted Distributed Graph Optimization (SW-DGO). D²RO formalizes a dimensionally weighted 5-component edge traversal cost function $C(u, v, t) = w_D D(u, v) + w_M W_{\text{mesh}}(u, v, t) + w_H H_{\text{prox}}(v, t) + w_R R_{\text{lock}}(u, v, t) + w_S S_{\text{trolley}}(v, t)$, unifying incremental heuristic graph repair (D* Lite), event-driven Vehicle-to-Vehicle (V2V) ad-hoc mesh telemetry with exponential decay, continuous 2D asymmetric Gaussian human proxemics, spatiotemporal directional corridor mutex locks, and non-holonomic vehicle safety clearance envelopes (S_{trolley}). We evaluate over 2,700 kinodynamic simulation trials in seven experiments: a comparative benchmark and component ablation in a retail supermarket, generalisation to clinical hospital (featuring Turnout Alcoves) and airport terminal topologies, two scalability sweeps, and two controlled mechanism experiments. D²RO completes 99.0% [94.6, 99.8] of missions in the primary supermarket benchmark at a median intimate-space exposure of 0 control ticks [IQR 0, 0], against 128 [123, 131] for socially blind Static A* and 204 [176, 276] for artificial potential fields (Wilcoxon signed-rank on $N = 100$ seed-paired trials, Holm-adjusted $p < 10^{-15}$). This social compliance is purchased with time rather than granted for free: D²RO's makespan of 47.18 ± 13.40 s is $2.6\times$ that of Static A* (18.00 ± 0.00 s), and we argue that this is the correct engineering trade for human-shared space. The unchanged planner generalizes across all three domains (99.0%, 100.0% and 95.0% success), and incremental D* Lite vertex repairs execute in 0.09–0.19 ms across crowd densities $N_{\text{humans}} \in [2..30]$ and fleet sizes $N_{\text{carts}} \in [2..12]$, under 0.4% of the 50 ms control-loop budget.

Index Terms—Multi-Agent Path Finding (MAPF), Socially-Aware Navigation, Distributed Graph Optimization, Vehicle-to-Vehicle (V2V) Mesh, Human Proxemics, D* Lite, Non-Holonomic Kinematics, Autonomous Mobile Robots (AMRs).

I. INTRODUCTION

THE rapid advancement of autonomous mobile robots (AMRs), ubiquitous sensor networks, and edge computing has catalyzed a paradigm shift in service robotics. While early automated guided vehicles (AGVs) operated primarily within segregated industrial warehouses—isolated from humans

behind physical safety barriers and governed by rigid guide-paths—the next generation of autonomous service fleets must operate directly within **human-shared, highly dynamic, and unstructured public environments** [1]–[3].

Prominent examples of such service fleets include:

- 1) **Smart Retail Shopping Trolleys (*Int-Cart*):** Autonomous motorized carts that escort shoppers, assist individuals with mobility impairments, execute in-store item collection for online grocery fulfillment, and autonomously return from checkout concourses to front-of-store cart depots [1].
- 2) **Clinical Hospital Pushchairs & Mobile Beds:** Robotic patient transport systems and mobile logistics couriers navigating high-stress hospital corridors between Emergency Departments (ER), Surgical Operating Suites (OR), Radiology/MRI suites, and Inpatient Wards.
- 3) **Airport Terminal Luggage Trolleys:** Heavy-duty autonomous baggage carts managing passenger luggage across open departure concourses, security screening bottlenecks, and narrow boarding gate piers.
- 4) **Smart Library & Micro-Fulfillment AGVs:** Material handling robots operating in narrow, single-file book stacks and retail storage aisles alongside human workers and patrons.

Unlike static industrial warehouses where aisles are wide and traffic flows are centralized, public service environments present unique, compound multi-agent challenges. In these settings, space is severely constrained by orthogonal shelving fixtures, corridors are frequently single-file, and human pedestrians move stochastically with complex browsing and loitering behaviors [2].

A. Core Research Motivations & Technical Bottlenecks

When deployed in complex, fixture-dense public service environments, traditional Multi-Agent Path Finding (MAPF) and reactive obstacle avoidance algorithms suffer from four fundamental technical bottlenecks:

- 1) *The Geometry-Kinematic Disconnect in Concave Fixtures (The ORCA Trap):* Classical continuous collision avoidance methods—such as Optimal Reciprocal Collision Avoidance (ORCA) [4] and Artificial Potential Fields (APF)—rely on local velocity-space half-planes or repulsive force vectors. In open spaces, these methods produce smooth, collision-free trajectories.

However, in realistic retail and hospital layouts filled with orthogonal 90° shelf corners, end-cap displays, and narrow U-bays, reactive methods are prone to **local potential minima**. When a robot encounters a pedestrian near a shelf corner, the repulsive force from the human and the repulsive vector from the orthogonal wall can cancel the goal attraction vector ($\|\mathbf{F}_{\text{net}}\| \rightarrow 0$), stalling the agent until the pedestrian moves on.

We stress that this is a *delay* mechanism, not necessarily a failure mechanism. Our measurements (Section VI-B) show that, given a time budget calibrated to the true mission duration, an artificial potential field completes every mission it is given; what it cannot do is complete them without repeatedly driving through people. The deficiency of reactive avoidance in these settings is therefore social and temporal rather than a simple inability to reach the goal, and the contribution of this paper should be read against that corrected baseline.

2) *Social Blindness of Static Shortest-Path Solvers (The A* Flaw)*: Traditional discrete graph planners (e.g., Static A*, Dijkstra) compute global paths based strictly on static Euclidean distance $D(u, v)$ [5]. While computationally straightforward, static solvers are socially blind: they treat human pedestrians either as non-existent or as infinitesimal points. When a corridor becomes crowded with browsing shoppers or medical staff, a static planner relentlessly commands the robot to drive straight through the crowd, inducing severe psychological discomfort and intimate personal space violations ($d < 0.8$ m).

3) *Symmetrical Live-Locks in Single-File Corridors*: In narrow architectural passages (e.g., supermarket grocery aisles, clinical hospital corridors) where corridor width W_{aisle} is less than twice the robot's safety envelope ($W_{\text{aisle}} < 2 \cdot r_{\text{safety}}$), bidirectional passing is geometrically impossible. When two opposing agents enter a single-file aisle simultaneously, reactive avoidance algorithms oscillate in place, producing permanent symmetrical live-locks [6].

4) *Information Isolation and Delayed Fleet Backtracking*: In decentralized multi-agent systems without communication, robots operate with myopic, line-of-sight sensing. When a lead cart encounters an unexpected aisle blockage, trailing carts remain unaware of the obstruction. Trailing units travel all the way to the blocked entrance before sensing the congestion locally, forcing complete reversals and causing severe inflation in fleet makespan [7].

B. Target Application Fields

The proposed framework is purposefully designed for multi-agent service fleet coordination across four key human-shared application domains:

- 1) **Smart Retail Commerce & Supermarkets**: Fleets negotiating narrow grocery aisles (Aisles 1–6), arterial promenades (**Action Alley**), end-cap displays, and cashier checkout queues.
- 2) **Clinical Healthcare Facilities & Hospitals**: Autonomous pushchairs enforcing strict priority hierarchies: emergency trauma vehicles maintain uninhibited right-of-way, while routine pushchairs yield dynamically into **Turnout Alcoves** (V_{alcove}).

- 3) **Aviation & Transportation Terminals**: Luggage trolleys navigating open concourses, security screening bottle-necks, and elongated boarding gate piers (Gates A1–A4, B1–B4).
- 4) **Public Facilities & Micro-Fulfillment Stacks**: AGVs operating in narrow book stacks and vertical shelving units where physical clearances are tight.

C. The Proposed Solution: The D²RO Framework

To overcome these bottlenecks, this paper proposes the **Distributed Dynamic Route Optimization (D²RO)** framework powered by **Socially-Weighted Distributed Graph Optimization (SW-DGO)**. Rather than treating navigation as isolated reactive avoidance or purely local search, D²RO establishes a **distributed, time-varying graph-cost field** $C_i(e, t)$ shared across peer agents:

$$C_i(e, t) = w_D \cdot C_{\text{geom}}(e) + w_M \cdot C_{\text{mesh}}(e, t) + w_H \cdot C_{\text{social}}(e, t) + w_R \cdot C_{\text{mutex}}(e, t) + w_S \cdot C_{\text{kinematic}}(i, e, t) \quad (1)$$

seamlessly unifying incremental heuristic graph repair (D* Lite) [8] with distributed anticipatory cost propagation, continuous Gaussian proxemics [2], [3], directional corridor reservation protocols, and non-holonomic vehicle clearance envelopes.

D. Research Aims & Key Contributions

The framework is designed against three explicit aims, and we return to each of them in the discussion and conclusion:

- A1 — **Socially-aware planning**: so that pedestrians' intimate space is not entered, accepting bounded additional travel time as the price. Evaluated in Section VI-C.
- A2 — **Distributed estimation of single-file conditions** using only peer-to-peer information within a bounded sensing radius. Evaluated in Section VI-G.
- A3 — **Architectural change detection and dissemination** of topologies within a 50 ms control budget. Evaluated in Sections VI-E and VI-F.

The primary scientific contributions of this research are:

- 1) **Distributed Anticipatory Edge-Cost Propagation & Horizon Extension**: We formalize a distributed dynamic graph-cost field in which localized physical and social perturbations observed by leading agents become peer-to-peer, time-decayed edge penalties $C_{\text{mesh}}(e, t)$. This extends each robot's effective planning horizon ($\mathcal{O}_i^{\text{effective}}(t)$) beyond its onboard line-of-sight sensing radius (R_s). Under a controlled experiment that isolates the mechanism (Section VI-G), agents reroute 10.70 ± 4.20 s earlier and more than halve their backtracking distance (1.08 ± 0.68 m vs. 2.73 ± 0.87 m).
- 2) **Distributed Directional Bottleneck Reservation Protocol**: We formalize an explicit distributed protocol for single-file topological bottlenecks ($W_{\text{corridor}} < 2r_{\text{safety}}$) via directional reservation tuples $\mathcal{L}_e = \langle \text{owner}, \vec{d}, t_{\text{acquire}}, t_{\text{expire}}, \text{priority} \rangle$ and Turnout Alcoves (V_{alcove}). In the controlled corridor experiment the protocol raises mission success from 36.0% to 88.0%. Our

instrumentation shows the mechanism operates by *cost-projected diversion* rather than by queueing, a distinction we analyse in Section VI-G.

- 3) **Kinodynamically & Socially Conditioned Incremental Optimization:** We unify asymmetric human proxemic discomfort fields (C_{social}) and vehicle-specific clearance envelopes ($C_{\text{kinematic}}$) directly into an incremental D* Lite graph repair engine, executing sub-millisecond updates (0.09–0.19 ms) without global re-heapification. Ablation shows the proxemic term is load-bearing: removing it collapses mission success from 100.0% to 11.0%.
- 4) **Comprehensive Multi-Domain Empirical Benchmark Suite:** We validate the framework across three distinct architectural topologies (Retail Supermarket, Clinical Hospital, and Airport Terminal) over 2,700 simulation trials in seven experiments. Every table and figure in this paper is generated programmatically from the committed raw data by a single analysis pipeline, with no hand-entered numbers; datasets and code are released under open licenses.

II. RELATED WORK & LITERATURE REVIEW

The proposed D²RO framework builds upon foundational work across Multi-Agent Path Finding (MAPF), dynamic graph replanning, local collision avoidance, ad-hoc mesh communication, physical trolley mechatronics, and human-aware navigation.

A. Decentralized and Lifelong MAPF

Stern *et al.* [5] define the classical MAPF problem and its variants, establishing the fundamental vertex and edge conflict models used in grid environments. To address continuous operations, Ma *et al.* [9] introduced Lifelong Multi-Agent Path Finding for Online Pickup and Delivery (MAPD), where agents dynamically receive tasks without resting.

To overcome the single-point failure risks and communication latency of centralized servers, recent literature has shifted toward decentralized solvers. Dergachev and Yakovlev [10] explored decentralized unlabeled MAPF using target and priority swapping, while Keskin *et al.* [11] presented a decentralized MAPF framework utilizing automated negotiation protocols. Furthermore, learning-based approaches have gained traction: Sartoretti *et al.* [12] introduced PRIMAL, utilizing reinforcement and imitation learning for decentralized pathfinding, and Skrynnik *et al.* [13] developed “Learn to Follow” to separate global heuristic sub-goal allocation from low-level local policies. Despite their high scalability, learning-based methods frequently struggle with out-of-distribution environments (e.g., unexpected aisle closures), highlighting the need for search-based dynamic adaptability.

B. Dynamic Replanning in Unknown and Stochastic Environments

In retail environments, the topological graph changes dynamically as aisles become blocked or crowded. Recalculating full paths from scratch for every agent is computationally

prohibitive. Koenig and Likhachev [8] addressed this with D* Lite, an incremental heuristic search algorithm that recalculates only the segments of a path affected by dynamic edge-cost changes.

The application of D* Lite to multi-agent systems was successfully demonstrated by Al-Mutib *et al.* [14], who utilized it for real-time path planning by treating peer paths as temporary obstacles. Additionally, Wagner and Choset [15] developed M*, dynamically varying the dimensionality of the search space only when agent paths conflict. D²RO adapts these incremental principles, allowing trolleys to dynamically inflate traversal costs based on real-time congestion data without recomputing the entire global map.

C. Kinematic Coordination and Local Collision Avoidance

While MAPF and D* Lite provide global waypoints, continuous kinematic control is required for safe micro-maneuvers when agents cross paths. Van den Berg *et al.* [4] introduced Optimal Reciprocal Collision Avoidance (ORCA), providing sufficient conditions for multiple robots to avoid collisions in continuous space without explicit communication.

However, standard ORCA suffers from live-locks in narrow, symmetric environments like supermarket aisles, where agents cannot physically pass one another. Dergachev and Yakovlev [6] specifically address this, proposing a system that uses continuous reciprocal collision avoidance but falls back on a locally confined MAPF instance when a deadlock is detected. D²RO leverages this exact synthesis: relying on continuous reactive models for open spaces, and spatiotemporal edge reservations for single-file corridors.

D. Multi-Robot Ad-Hoc Communication Protocols

The transition from centralized control to a truly distributed D²RO system requires robust peer-to-peer communication. Gielis *et al.* [7] emphasize the critical need for co-designing robotic planning algorithms alongside network constraints.

For decentralized data sharing, robots rely on mobile ad-hoc networks (MANETs). Slyusar and Kulich [16] evaluated routing protocols for MANETs in multi-robot exploration. Additionally, Edwige *et al.* [17] investigated robot communication within Swarm SLAM, demonstrating how independent agents merge local spatial data over a distributed mesh. In D²RO, this translates to an event-driven telemetry protocol where agents broadcast localized edge-cost penalties across a V2V mesh, allowing distant agents to proactively reroute.

E. Indoor Positioning and Physical Hardware Implementation

Unlike simulated grids, physical shopping trolleys require absolute spatial grounding and customized mechatronics. Zafari *et al.* [18] provide a comprehensive survey of indoor localization technologies, highlighting the superiority of Ultra-Wideband (UWB) and BLE for centimeter-level accuracy in GPS-denied environments. Clark *et al.* [19] expanded on this with the TEAM framework, demonstrating effective trilateration and mapping utilizing a localized robotic network, while Nugraha *et al.* [20] proved that fusing Indoor Positioning Systems

(IPS) with wheel odometry via Extended Kalman Filters (EKF) drastically reduces navigation drift.

Bringing these concepts into physical retail spaces, Mohamad Azlan *et al.* [1] developed the *Int-Cart*, an autonomous mobile trolley robot. Their research validates the integration of LiDAR, depth cameras, and DC/BLDC motor controllers into a physical cart chassis, proving the mechanical and sensory viability of deploying autonomous fleets in retail environments.

F. Human-Aware Navigation & Research Gap

Recent benchmark frameworks, such as HA-VLN 2.0 [21], emphasize that robots cannot treat humans simply as “moving cylindrical obstacles.” Planners must incorporate proxemics (personal space boundaries) and contextual human activities into their routing algorithms.

The Research Gap: There remains a distinct lack of hybrid frameworks that fuse proactive, mesh-informed global graph updates with reactive human-aware collision avoidance in highly constrained physical retail spaces. Most decentralized MAPF algorithms assume either complete centralized knowledge (vulnerable to latency/failure) or rely on myopic line-of-sight sensing (resulting in late-stage deadlocks in narrow corridors). The D²RO framework bridges this gap.

III. PROBLEM FORMULATION & THE SW-DGO FRAMEWORK

A. Topological Roadmap and Kinematic Model

Let the indoor facility floorplan be represented by an undirected planar topological graph $G = (V, E)$, where $V = \{v_1, v_2, \dots, v_n\}$ denotes navigable topological waypoints (e.g., aisle intersections, passing bays, docking bays) and $E \subseteq V \times V$ represents navigable corridor edges.

Each autonomous service cart $i \in \{1, \dots, N_{\text{carts}}\}$ is governed by non-holonomic unicycle kinematics:

$$\begin{aligned}\dot{x}_i(t) &= v_i(t) \cos \theta_i(t), \\ \dot{y}_i(t) &= v_i(t) \sin \theta_i(t), \quad \dot{\theta}_i(t) = \omega_i(t)\end{aligned}\quad (2)$$

subject to physical actuator bounds $|v_i(t)| \leq v_{\max}$ and $|\omega_i(t)| \leq \omega_{\max}$.

B. The Distributed Dynamic Edge-Cost Field Formulation

Rather than relying on centralized graph allocation or isolated myopic sensing, SW-DGO establishes a **distributed, time-varying edge-cost field** $C_i(u, v, t)$ maintained locally by each autonomous agent i . The edge traversal cost $e = (u, v) \in E$ is formalized as:

$$\begin{aligned}C_i(u, v, t) &= w_D \cdot C_{\text{geom}}(u, v) + w_M \cdot C_{\text{mesh}}(u, v, t) \\ &\quad + w_H \cdot C_{\text{social}}(v, t) + w_R \cdot C_{\text{mutex}}(u, v, t) \\ &\quad + w_S \cdot C_{\text{kinematic}}(i, v, t)\end{aligned}\quad (3)$$

where the calibrated dimensionless weights $\mathbf{w} = [w_D, w_M, w_H, w_R, w_S] = [1.0, 1.5, 2.0, 1.0, 1.2]$ balance physical metric progress against social proxemics, collaborative telemetry, and vehicle clearance:

- $w_D = 1.0$: Baseline Euclidean physical distance (C_{geom} in meters).
- $w_M = 1.5$: Weights collaborative V2V congestion alerts $1.5\times$ higher than metric distance, ensuring trailing carts proactively reroute through clear parallel corridors rather than deadheading into blocked aisles.
- $w_H = 2.0$: Prioritizes pedestrian personal-space comfort $2.0\times$ above travel distance, enforcing compliant yielding and wide social berths.
- $w_R = 1.0$: Directional corridor mutex multiplier ($C_{\text{mutex}} = \infty$ during active contention, 0 otherwise), which renders a contested corridor untraversable for the deferring agent in its own cost field. We claim this empirically, not as a proven exclusivity guarantee.
- $w_S = 1.2$: Kinetic clearance penalty preventing chassis corner scrapes ($1.2\times$ baseline) and anti-tailgating violations.

1) 1. Intrinsic Metric Geometry $C_{\text{geom}}(u, v)$:

$$C_{\text{geom}}(u, v) = \|\mathbf{p}_u - \mathbf{p}_v\|_2 + \alpha_{\text{turn}} |\Delta\theta| \quad (4)$$

where α_{turn} penalizes sharp angular steering deviations to favor smooth corridor alignment.

2) 2. Distributed Anticipatory Cost Propagation & Horizon Extension $C_{\text{mesh}}(u, v, t)$: In decentralized multi-agent systems, myopic line-of-sight sensing restricts an agent's local obstacle horizon to $\mathcal{O}_i^{\text{local}}(t) = \{\mathbf{x} \mid \|\mathbf{x} - \mathbf{p}_i(t)\| \leq R_s\}$. To eliminate the severe makespan penalty of deadheading into remote blockages, SW-DGO formalizes an effective perception horizon extension:

$$\mathcal{O}_i^{\text{effective}}(t) = \mathcal{O}_i^{\text{local}}(t) \cup \bigcup_{j \in \mathcal{N}_i} \mathcal{O}_j(t - \tau_{ij}) \quad (5)$$

When agent j discovers a congestion event at edge (u, v) at timestamp t_k , it broadcasts a localized event-driven telemetry packet over the ad-hoc mesh with bounded time-to-live (TTL = 3). The received penalty decays exponentially over time:

$$C_{\text{mesh}}(u, v, t) = \sum_{k \in \mathcal{M}(u, v, t)} \gamma_k \cdot \exp(-\lambda_{\text{decay}}(t - t_k)) \quad (6)$$

where γ_k is the initial congestion penalty amplitude and λ_{decay} is the temporal decay rate, guaranteeing that stale congestion information naturally attenuates without permanently poisoning the routing topology.

3) 3. Continuous 2D Asymmetric Human Proxemics $C_{\text{social}}(v, t)$: To model human psychological personal space (Hall's Proxemics), we formulate an asymmetric directional Gaussian potential field aligned with pedestrian heading θ_h :

$$H_{\text{prox}}(\mathbf{p}, \mathbf{h}_j, \theta_h) = A_j \cdot \exp\left(-\frac{1}{2} \left[\left(\frac{x_{\text{local}}}{\sigma_x}\right)^2 + \left(\frac{y_{\text{local}}}{\sigma_y}\right)^2 \right]\right) \quad (7)$$

where relative coordinates are transformed into the pedestrian's local body frame:

$$\begin{bmatrix} x_{\text{local}} \\ y_{\text{local}} \end{bmatrix} = \begin{bmatrix} \cos \theta_h & \sin \theta_h \\ -\sin \theta_h & \cos \theta_h \end{bmatrix} \begin{bmatrix} x - x_j \\ y - y_j \end{bmatrix} \quad (8)$$

TABLE I: System Simulation Parameters & Physical SI Metric Calibration.

Parameter Description	Symbol	Physical Value (SI)
Spatial Resolution Scale	s_{res}	1 px = 0.03 m (1 m = 33.33 px)
Physics Integration Rate	Δt	0.05 s (20 Hz control loop)
GUI Rendering Rate	f_{gui}	60 FPS (16.6 ms interpolation)
Chassis Dimensions	$L \times W$	0.72 m $\times$ 0.48 m (24 $\times$ 16 px)
Inscribed Safety Radius	r_{robot}	0.40 m (13.3 px)
Maximum Linear Speed	v_{max}	1.20 m/s (40.0 px/s)
Maximum Angular Rate	ω_{max}	2.50 rad/s
Shelf Clearance Margin	d_{shelf}	0.54 m (18.0 px)
Following Gap Distance	d_{follow}	1.08 m (36.0 px)
V2V Comm Range	R_{mesh}	10.50 m (350.0 px)
V2V Packet Size	S_{pkt}	64 Bytes
Mesh Hop Limit	TTL	3 hops
Proxemic Front Space	σ_{front}	1.35 m (45.0 px)
Proxemic Lateral Space	σ_{side}	0.90 m (30.0 px)
Proxemic Rear Space	σ_{rear}	0.60 m (20.0 px)
Onboard Sensing Radius	R_s	7.20 m (240.0 px)
Arrival Tolerance (all planners)	r_{arrive}	0.84 m (28.0 px)

with asymmetric longitudinal variance reflecting expanded front personal space:

$$\sigma_x = \begin{cases} 1.35 \text{ m}, & x_{\text{local}} \geq 0 \text{ (front)} \\ 0.60 \text{ m}, & x_{\text{local}} < 0 \text{ (rear)} \end{cases} \quad \sigma_y = 0.90 \text{ m} \quad (9)$$

The continuous social edge penalty is integrated along the edge trajectory segment:

$$C_{\text{social}}(u, v, t) = \int_0^1 H_{\text{prox}}((1 - \tau)\mathbf{p}_u + \tau\mathbf{p}_v, t) d\tau \quad (10)$$

4) *4. Distributed Directional Bottleneck Reservation Protocol* $C_{\text{mutex}}(u, v, t)$: In single-file architectural passages where aisle width cannot accommodate bidirectional passing ($W_{\text{corridor}} < 2r_{\text{safety}}$), mutual exclusion is coordinated via a directional reservation tuple:

$$\mathcal{L}_e(t) = \langle \text{AgentID}, \vec{d}_{\text{traversal}}, t_{\text{acquire}}, t_{\text{expire}}, \text{Priority} \rangle \quad (11)$$

When agent i claims edge (u, v) , the reverse edge cost for opposing agents is set to:

$$C_{\text{mutex}}(v, u, t) = \begin{cases} \infty, & (u, v) \text{ reserved by peer } j \neq i \\ 0, & \text{otherwise} \end{cases} \quad (12)$$

An opposing agent observes this infinite cost and, because the reservation enters its *graph* rather than its controller, D* Lite reroutes it through a parallel aisle or a Turnout Alcove (V_{alcove}) instead of halting at the corridor mouth. Section VI-G shows empirically that this diversion—rather than queueing—is the operative behaviour.

5) *5. Kinetic Chassis Safety Clearance Envelope* $C_{\text{kinematic}}(i, v, t)$: To prevent chassis scrapes against orthogonal shelf corners and enforce dynamic following gaps:

$$C_{\text{kinematic}}(i, \mathbf{p}, t) = \sum_{j \neq i} A_{\text{trolley}} \cdot \exp\left(-\frac{\|\mathbf{p} - \mathbf{p}_j(t)\|^2}{2\sigma_{\text{trolley}}^2}\right) \quad (13)$$

enforcing an 18 px (0.54 m) shelf clearance buffer and a 36 px (1.08 m) anti-tailgating following distance.

Algorithm 1 Socially-Weighted Distributed Graph Optimization (SW-DGO)

Require: Roadmap graph $G = (V, E)$, Start s_{start} , Goal s_{goal} , V2V Mesh $\mathcal{N}$

- 1: Initialize local D* Lite planner with G and s_{goal}
- 2: Compute initial shortest path $\Pi \leftarrow \text{ComputeShortestPath}()$
- 3: **while** $s_{\text{current}} \neq s_{\text{goal}}$ **do**
- 4: $\Delta\mathcal{M} \leftarrow \text{FetchInboundMeshPackets}(\mathcal{N})$
- 5: cost_changed $\leftarrow$ **false**
- 6: **for** each packet $\text{pkt} \in \Delta\mathcal{M}$ **do**
- 7: $(u, v) \leftarrow \text{pkt.edge}$
- 8: Update edge cost: $c(u, v) \leftarrow C(u, v, t)$ via Eq. (3)
- 9: D*Lite.NotifyEdgeCostChange(u, v)
- 10: cost_changed $\leftarrow$ **true**
- 11: **end for**
- 12: Sample human positions $\{h_j\}$ within radius $R_{\text{sense}} = 240$ px
- 13: **if** ProxemicsChanged($\{h_j\}$) **then**
- 14: Update $H_{\text{prox}}(u, v, t)$ for local edges
- 15: cost_changed $\leftarrow$ **true**
- 16: **end if**
- 17: **if** cost_changed **then**
- 18: D*Lite.ComputeShortestPath()
- 19: $\Pi \leftarrow \text{D*Lite.ExtractPath}()$
- 20: **end if**
- 21: $s_{\text{next}} \leftarrow \text{GetNextWaypoint}(\Pi)$
- 22: **if** EdgeIsSingleFile($s_{\text{current}}, s_{\text{next}}$) **then**
- 23: AcquireCorridorLock($s_{\text{current}}, s_{\text{next}}$)
- 24: Broadcast(LOCK_REQUEST($s_{\text{current}}, s_{\text{next}}$))
- 25: **end if**
- 26: Execute non-holonomic step toward s_{next} with S_{trolley} envelope
- 27: Decay active mesh penalties: $W_{\text{mesh}} \leftarrow W_{\text{mesh}} \cdot e^{-\lambda\Delta t}$
- 28: **end while**
- 29: **return** Mission Complete

C. Comparison with Alternative Simulation Paradigms

To contextualize the architectural design of D²RO, we contrast it with four standard robotics simulation paradigms:

- 1) **Classical 2D Grid-World MAPF (CBS, PBS, EECBS):** Assumes synchronous discrete time steps and point holonomic agents. Ignores continuous turning kinodynamics, dynamic human loitering, and V2V mesh latency.
- 2) **Pure Continuous Multi-Agent Navigators (ORCA, Social Force):** Operates in continuous velocity space without topological roadmaps. Lacks global topological awareness, so an agent driven into an infeasible velocity state in a narrow corridor has no route-level recourse. Our implementation exhibits this mode (Section VI-B), though we do not generalise from it to the published algorithm.
- 3) **Network-Centric Packet Simulators (NS-3, OM-NeT++):** Models high-fidelity RF propagation and MAC collisions, but lacks robot kinodynamic physics and graph search.
- 4) **3D Rigid-Body Physics Engines (Gazebo, Isaac Sim):**




Fig. 1: Supermarket environment floorplan with SW-DGO planned trajectories, aisle shelves, Action Alley promenade, human Gaussian halos, and Cart Depots.




Fig. 2: Hospital autonomous pushchair floorplan featuring Emergency Trauma (ER), Sterile OR/MRI, Clinical Wards, and Turnout Alcoves for dynamic yielding.

Models full multi-body dynamics, contact mechanics, and sensor ray-tracing. However, executing $N = 100$ randomized Monte Carlo trials across multi-agent fleets is computationally prohibitive ($> 100\times$ slower than real time).

- 5) **The SW-DGO Hybrid Approach (The Goldilocks Paradigm):** Couples non-holonomic unicycle kinodynamics ($dt = 0.05$ s, 60 FPS) with continuous Gaussian proxemics, an ad-hoc V2V mesh simulator, and incremental D^{*} Lite graph repair, providing deterministic, real-time evaluation (< 0.15 ms).

IV. MULTI-DOMAIN SIMULATION TOPOLOGIES

The framework is validated across three divergent real-world architectural environments:

A. Retail Supermarket (Int-Cart Fleet)

As shown in Fig. 1, the supermarket environment features 6 parallel grocery aisles ($W_{\text{aisle}} = 70$ px), transverse cross-promenades (Back Promenade, Action Alley, Front Concourse), cashier registers, and front cart return depots.

B. Clinical Hospital (Pushchair Fleet)

As shown in Fig. 2, the hospital floorplan includes Emergency Trauma (ER), Sterile Operating Theatres (OR/MRI), Inpatient Wards, and designated **Turnout Alcoves** (V_{alcove}) along single-file corridors to grant priority clearance to emergency vehicles.

C. Airport Terminal (Luggage Trolleys)

As shown in Fig. 3, the airport concourse combines a large open-plan departure plaza with narrow security screening chokepoints and boarding gate piers (Gates A1–A4, Gates B1–B4).




Fig. 3: Airport terminal autonomous luggage trolley concourse simulation with Check-in Banks, Security screening, open plaza, and Gate Piers.

V. SPATIAL PROXEMIC HEATMAPS & SPATIOTEMPORAL ANALYSIS

A. Social Detour Heatmap Analysis

Fig. 4 maps the continuous 2D Gaussian discomfort field $H_{\text{prox}}(x, y)$ alongside executed trajectories. Static A^* (dashed red) charges directly through the dense crowd in Aisle 3, whereas D²RO (solid blue) executes a proactive detour along Action Alley and Aisle 2, achieving zero intimate personal space violations.

B. Spatiotemporal Turnout Alcove Resolution

Fig. 5 illustrates the longitudinal space-time trajectory (X, t) in clinical corridors. Emergency Pushchair P_1 maintains full velocity under priority lock ($R_{\text{lock}} = \infty$), while routine Pushchair P_2 pulls into the Turnout Alcove at $X = 650$ px ($t = 4.5$ s $\rightarrow$ 10.5 s), resolving the conflict without either agent reversing. The figure illustrates intended behaviour; the quantitative evidence is in Section VI-G.




Fig. 4: Continuous 2D Gaussian Discomfort Field $H_{\text{prox}}(x, y)$ heatmap and trajectory overlay comparing Static A^* (blind path through Aisle 3 crowd) against D²RO proactive social detour along Action Alley.




Fig. 5: Spatiotemporal time-space trajectory diagram of Turnout Alcove resolution: Emergency Pushchair P_1 maintains full velocity under priority lock $R_{\text{lock}} = \infty$, while routine Pushchair P_2 yields inside the alcove bay.

C. Airport Concourse Flow Streamlines

Fig. 6 demonstrates 2D velocity streamlines around high-density passenger clusters in open airport concourses, highlighting smooth non-holonomic curvature without erratic stopping.

VI. EXPERIMENTAL RESULTS & DISCUSSION

A. Experimental Setup & Randomization Protocol

All experiments were conducted using the decoupled 2D kinodynamic simulation framework integrating non-holonomic unicycle dynamics ($\Delta t = 0.05$ s, 20 Hz control loop) across 1200×800 px (36.0 m $\times$ 24.0 m) floorplans. To ensure statistical rigor and reproducibility:

- **Random Seed Policy:** Each evaluation executed $N = 100$ independent Monte Carlo trials seeded deterministically. Crucially, every planner in a given trial receives the *same* seed, so comparisons are **seed-paired** rather than independent: the pedestrian crowd a baseline faces is the crowd D²RO faces.




Fig. 6: Airport open concourse vector flow streamlines and multi-agent luggage cart trajectories smoothly navigating around high-density passenger clusters.

- **Dynamic Pedestrian Behavior:** Pedestrians move with stochastic velocities $v_h \in [0.8, 1.2]$ m/s, random way-point heading deflections, and a 30% browsing loitering probability ($p_{\text{loiter}} = 0.30$) generating dynamic Gaussian discomfort Halos.
- **Common Completion Criterion:** Within any given experiment all planners share an identical arrival tolerance of 0.84 m (approximately one cart length) and an identical time budget. An earlier revision applied planner-specific tolerances, scoring the proposed method against a $2.3\times$ more lenient criterion than its baselines; that asymmetry is removed and every result reported here uses the common criterion. Budgets are set well above the longest observed mission for the experiment class so that a timeout indicates genuine non-completion rather than a truncated run: $T_{\text{max}} = 180$ s for the benchmark, ablation, cross-domain and crowd-density experiments, 240 s for fleet-size scaling (where queueing grows with fleet size), and 120 s for the two-cart mechanism experiments.
- **Statistical Analysis:** Because the trials are seed-paired and several outcome distributions are zero-inflated and strongly right-skewed, we do not assume normality. Continuous outcomes are compared with the **Wilcoxon signed-rank** test on within-pair differences, supported by paired bootstrap confidence intervals; binary mission success is compared with **McNemar's exact** test. p -values are **Holm-adjusted** within each family of comparisons, and every reported p belongs to one specific outcome—no single significance claim is attached to several unrelated metrics. Skewed outcomes are summarised as median [IQR]; symmetric ones as $\mu \pm \sigma$.

B. Physical Insights and Failure Modes Analysis

1) **Artificial Potential Fields: a social failure, not a navigational one:** An earlier revision of this study reported that APF achieves 0.0% mission success and built a local-minima argument upon it. Under the common completion criterion and a time budget calibrated to true mission durations, that result does not reproduce: **APF completes 100% of its missions at**

TABLE II: Comparative benchmark in the retail supermarket domain ($N = 100$ paired Monte Carlo trials, identical seeds across all planners). Social exposure is reported as median [IQR] because the distribution is zero-inflated and strongly right-skewed. p -values are Holm-adjusted within the family of comparisons and each refers to the single outcome on its own row.

Algorithm	Success (95% CI)	Makespan, successful (s)	Deadlocks	Intimate exposure, median [IQR]	p (exposure)
D ² RO (proposed)	99.0% [94.6, 99.8]	47.18 ± 13.40	0.00 ± 0.00	0 [0, 0]	—
Static A*	100.0% [96.3, 100.0]	18.00 ± 0.00	0.00 ± 0.00	128 [123, 131]	$5.1e - 16$
APF	100.0% [96.3, 100.0]	34.54 ± 0.16	0.00 ± 0.00	204 [176, 276]	$2.7e - 16$
Reactive ORCA	0.0% [0.0, 3.7]	no success	0.00 ± 0.00	229 [0, 766]	$1.5e - 13$
Decentralized Local MAPF	0.0% [0.0, 3.7]	no success	0.00 ± 0.00	128 [123, 132]	$5.1e - 16$




Fig. 7: Comparative evaluation across (a) mission success rate, (b) makespan of successful missions, and (c) intimate-space exposure (median, IQR whiskers). Panels (b) and (c) show the trade-off directly: D²RO is slower than Static A* and APF while intruding on personal space substantially less. Success rates in (a) are statistically indistinguishable between D²RO, Static A* and APF.

34.54 ± 0.16 s. The previous figure was an artefact of a 35.0 s timeout inherited from an earlier, incorrectly scaled kinematic model, not a property of the algorithm. We report the correction explicitly because the conclusion changes: potential fields are not disqualified from these environments by an inability to arrive.

What disqualifies them is *how* they arrive. APF records the highest intimate-space exposure of any planner evaluated—median 204 ticks [IQR 176, 276], against 0 [0, 0] for D²RO (Wilcoxon signed-rank, $\Delta = -220.7$ [−239.5, −201.4], Holm-adjusted $p = 2.7 \times 10^{-16}$). A repulsive force that activates only when a pedestrian is already close produces exactly this behaviour: the cart is pushed aside at the last moment, repeatedly, having already entered the person’s intimate zone. The deficiency is one of *anticipation*, which is precisely what a cost field evaluated over a planning horizon supplies and a local force law cannot.

2) *ORCA and Decentralized Local MAPF: results we decline to over-interpret*: Our implementations of Reactive ORCA and Decentralized Local MAPF complete no missions under the common arrival criterion. Instrumented traces show ORCA agents stalled for 91–99% of control ticks, having covered only a fraction of a ~ 17 m mission, which is consistent with the velocity-obstacle infeasibility mode described in Section I-A1: in a corridor narrower than twice the safety envelope, the half-plane constraints admit no feasible velocity and the agent halts.

We deliberately stop short of the stronger claim. ORCA is a widely deployed algorithm, and an in-house implementation returning 0% where the published method is used successfully in practice is at least as likely to indicate a defect in our

implementation as a fundamental limitation of the method. **We therefore report these as properties of our implementations, not of the algorithms**, and we identify validation against a reference implementation on a canonical benchmark as required work before any general claim about ORCA’s behaviour in single-file corridors can be sustained. The same caution applies to Local MAPF, where traced trials show agents finishing within ~ 1.3 m of their goals rather than livelocked—a near-miss, not the permanent token-swapping deadlock previously asserted. Consequently, none of the contributions claimed in this paper rest on the ORCA or MAPF baselines; they rest on the comparison against Static A* and APF, both of which complete their missions.

3) *Anticipatory Horizon Extension vs. Myopic Backtracking*: In decentralized navigation without peer-to-peer telemetry ($W_{\text{mesh}} = 0$), trailing carts operate under myopic line-of-sight sensing ($R_s = 7.2$ m). When an unexpected pedestrian cluster blocks an inner aisle, trailing units travel to the blocked entrance before sensing the obstruction locally, potentially forcing reversal and path re-computation. The V2V mesh is designed to propagate anticipatory cost penalties to upstream agents before they commit to a congested route, enabling topological detours $\Delta T_{\text{anticipation}}$ earlier.

Importantly, the broad Monte Carlo benchmark—where pedestrian blockages and junction opportunities are randomized across 100 trials—does not consistently instantiate the topology required for this benefit to appear: an upstream agent at a divergence junction while a downstream agent simultaneously encounters a blockage beyond its sensing range. In those trials the mesh pays its overhead (cost-field propagation and broadcast-triggered replanning) without collecting its return, which is why the no-mesh ablation completes the benchmark faster than the full system (38.23 ± 16.88 s vs. 47.38 ± 13.96 s; Table III). This does not invalidate the mechanism, but neither does the benchmark evidence it: a general-scenario experiment is simply not a test of anticipatory communication. Section VI-G therefore evaluates the mesh under conditions that do instantiate the required topology.

C. Comparative Benchmark Analysis: an explicit trade-off

Table II and Fig. 7 do not show the proposed method dominating on every axis, and we present them as what they are—a trade-off that we argue is favourable for human-shared space.

Where D²RO leads. On social compliance the separation is large and consistent. D²RO records a median intimate-space exposure of 0 ticks [IQR 0, 0]: in the majority of trials the

fleet never enters a pedestrian’s intimate zone at all. Static A^* records 128 [123, 131] and APF 204 [176, 276]. We report medians because the distribution is zero-inflated and strongly right-skewed—D²RO’s own mean is inflated by a small number of trials in which a pedestrian steps into an already-committed corridor—and a $\mu \pm \sigma$ summary would imply a symmetry that does not exist.

What that compliance costs. D²RO completes missions in 47.18 ± 13.40 s against Static A^* ’s 18.00 ± 0.00 s: a $2.6\times$ makespan penalty (Wilcoxon signed-rank, $\Delta = 30.51$ s, $d_z = 1.62$, Holm-adjusted $p = 3.8 \times 10^{-17}$). It also does not lead on raw success: 99.0% [94.6, 99.8] against 100.0% for both Static A^* and APF, a difference of a single trial that McNemar’s exact test does not distinguish from chance ($p = 1$). Claiming a success-rate advantage over these two baselines would not be supportable, and we do not make one.

Why the trade is the right one. A cart that reaches the depot four seconds sooner by driving through a queue of shoppers has not solved the problem this paper poses. The design objective (Section I-D) is socially compliant routing subject to completing the mission, and against that objective the measured outcome is the intended one: near-total elimination of personal-space intrusion, purchased with time, at no measurable cost in reliability.

D. Component Ablation Analysis

The ablation in Table III distinguishes the terms of Eq. (3) that are load-bearing in this benchmark from those that are not, and we report both outcomes.

The proxemic term carries the framework. Removing H_{prox} collapses mission success from 100.0% to 11.0% and raises the discomfort integral from 0.05 ± 0.13 to 13.14 ± 3.57 , with makespan degrading from 47.38 ± 13.96 s to 172.05 ± 26.67 s. Without a social cost the planner routes straight into crowds, then relies on the reactive yielding layer to extricate itself, which it frequently cannot do within the time budget. This single term accounts for most of what distinguishes D²RO from Static A^* .

The mesh and mutex terms are not exercised by this benchmark. Removing W_{mesh} or R_{lock} leaves success at 100.0% and *reduces* makespan, from 47.38 ± 13.96 s to 38.23 ± 16.88 s and 37.86 ± 15.13 s respectively. Read alone, this ablation argues against both mechanisms. The explanation is that a randomized supermarket scenario rarely instantiates the topology each mechanism addresses—a downstream blockage outside a trailing agent’s sensing range, or two agents committed to the same single-file corridor from opposite ends—while their costs (broadcast-triggered replanning, conservative corridor deferral) are paid in every trial. Section VI-G therefore evaluates both under controlled conditions that do instantiate those topologies. We consider it important to report the negative general-scenario result alongside the positive controlled one rather than only the latter.

The safety envelope shows a weaker effect than previously claimed. Removing S_{trolley} raises shelf-corner scrapes from 193.05 ± 169.90 to 271.71 ± 84.16 . This is a real degradation, but it is a 41% increase rather than the elimination of scraping:

the envelope reduces fixture contact, it does not prevent it, and the full configuration itself records substantial contact. An earlier revision reported 0.00 scrapes for the full system against 77.72 ± 11.94 without it; that contrast was an artefact of the superseded metric semantics and does not survive measurement under a single shared threshold. The high variance in both arms further suggests scrape counts are dominated by a few tight-clearance geometries rather than distributed evenly across trials.

E. Cross-Domain Failure Mode Analysis

Aim A3 requires that the planner transfer across dissimilar architectures without retuning. Table IV supports this: with identical weights and no domain-specific parameters, D²RO completes 99.0% of missions in the retail supermarket, 100.0% in the clinical hospital and 95.0% in the airport terminal. Generality is therefore demonstrated; what differs across domains is not reliability but the *character* of the routing effort, which the remaining columns expose.

The hospital is the easiest domain, not the hardest. It records the highest success rate and the lowest replan count (342.47 ± 64.94), despite having the narrowest corridors. Its topology is strongly channelised: routes are largely predetermined by the corridor structure, so there are few alternatives to evaluate and little for the planner to reconsider once committed. Its longer mean transit (31.78 ± 3.89 s) reflects distance and mandatory alcove yields, not indecision.

The airport is the hardest, and for an instructive reason. It requires 959.2 ± 359.8 replans—nearly twice the supermarket’s 487.9 ± 151.1 —and is the only domain with non-zero median intimate exposure (62 [18, 130], against 0 [0, 0] in both other domains). An open concourse offers a large, continuously changing set of admissible routes, so a crowd shifting a few metres alters the cost field enough to justify a new plan; the planner spends its effort re-deciding rather than executing. The wide interquartile range indicates the exposure is concentrated in trials where the fleet must cross a dense gate-area crowd, rather than being incurred uniformly. Open topologies thus stress the framework differently from constrained ones—the difficulty is choice abundance, not scarcity—which motivates crowd-density-adaptive scheduling of w_H and w_M , and a replan-hysteresis term, as future work.

We note that these figures supersede an earlier revision reporting 80.0% airport success. In that revision the airport scenario reset its random seed to a constant, so its 100 nominal trials were 100 repetitions of a single trial; the defect is corrected and the domain now varies per trial as the other two do.

F. Decoupled Scalability Analysis

We vary crowd density and fleet size independently ($N = 100$ trials per point) so that the two effects are not confounded. The two dimensions behave very differently, and only one of them is benign.

Crowd density scales benignly (Fig. 8). Increasing pedestrians from 2 to 30 raises replan latency only from 0.09 ± 0.01 ms to 0.17 ± 0.05 ms and makespan from 30.10 s to 56.30 s, while

TABLE III: Component ablation of the five cost terms in the retail supermarket domain ($N = 100$ trials per configuration). Each row removes exactly one term from Eq. (1); all other parameters are held fixed.

Configuration	Success (95% CI)	Makespan (s)	Discomfort	Deadlocks	Shelf scrapes
Full D²RO	100.0% [96.3, 100.0]	47.38 $\pm$ 13.96	0.05 $\pm$ 0.13	0.00 $\pm$ 0.00	193.05 $\pm$ 169.90
w/o V2V mesh W_{mesh}	100.0% [96.3, 100.0]	38.23 $\pm$ 16.88	0.09 $\pm$ 0.43	0.00 $\pm$ 0.00	191.59 $\pm$ 182.39
w/o corridor mutex R_{lock}	100.0% [96.3, 100.0]	37.86 $\pm$ 15.13	0.09 $\pm$ 0.38	0.00 $\pm$ 0.00	186.13 $\pm$ 164.45
w/o proxemics H_{prox}	11.0% [6.3, 18.6]	172.05 $\pm$ 26.67	13.14 $\pm$ 3.57	0.00 $\pm$ 0.00	244.86 $\pm$ 154.18
w/o safety envelope S_{trolley}	100.0% [96.3, 100.0]	43.02 $\pm$ 8.84	0.05 $\pm$ 0.17	0.00 $\pm$ 0.00	271.71 $\pm$ 84.16

TABLE IV: Cross-domain generalisation of the unchanged planner across three topologically distinct environments ($N = 100$ trials per domain). Intimate exposure is reported as median [IQR]; the distribution is zero-inflated.

Domain	Success (95% CI)	Makespan (s)	Mean transit (s)	Intimate exposure	D* Lite replans
Retail Supermarket	99.0% [94.6, 99.8]	49.89 $\pm$ 19.94	25.90 $\pm$ 5.57	0 [0, 0]	487.9 $\pm$ 151.1
Clinical Hospital	100.0% [96.3, 100.0]	46.12 $\pm$ 10.12	31.78 $\pm$ 3.89	0 [0, 0]	342.5 $\pm$ 64.9
Airport Terminal	95.0% [88.8, 97.8]	74.63 $\pm$ 35.71	36.84 $\pm$ 15.01	62 [18, 130]	959.2 $\pm$ 359.8




Fig. 8: Crowd-density scalability at a fixed fleet of four carts. Error bars are ± 1 SD, except (c) which shows the Wilson 95% interval.

success falls just 100.0% $\rightarrow$ 98.0%. Incremental repair is doing its job: a denser crowd perturbs more edges, but D* Lite repairs only the affected subgraph, so cost grows with the perturbation rather than with the map. Even at the highest density the planner consumes under 0.4% of the 50 ms control budget, satisfying the real-time half of aim A3 with three orders of magnitude of headroom.

Fleet size does not (Fig. 9), and this is the framework’s clearest limitation. Success holds near 100% up to 8 carts, then falls to 85.0% at 10 and 78.0% at 12, while makespan roughly doubles (50.99 s $\rightarrow$ 110.17 s). The mechanism is visible in the mesh column: broadcast volume grows super-linearly, from 8.50 ± 5.72 packets at two carts to 1158.60 ± 336.83 at twelve, because every agent broadcasts observations that every other agent must incorporate. Beyond roughly eight carts in a 36×24 m floorplan the fleet begins to congest itself: agents increasingly replan around each other’s reservations and telemetry rather than around pedestrians.

We state this bound plainly rather than describing the degradation as graceful. For the retail and clinical deployments motivating this work—where a realistic fleet in a single zone is a handful of units—the operating regime sits inside the flat part

of the curve. But the present design should not be represented as scaling to large fleets in a confined space, and mesh-traffic suppression (relevance filtering, spatial scoping of broadcasts) is the natural next step for that regime.

G. Mechanism-Specific Validation

Because the broad Monte Carlo benchmark does not reliably trigger the specific topological conditions under which V2V mesh anticipation and corridor mutex locking provide their primary contribution, we designed two controlled experiments that directly exercise each mechanism. These complement the general-scenario evaluation and provide targeted evidence for the core algorithmic novelties.

1) *Experiment A: V2V Mesh Anticipation ($N = 50$ trials):* Two carts are positioned such that Cart A is in front and Cart B is 10+ m behind (outside Cart A’s sensing range). A pedestrian cluster blocks Cart A’s forward aisle mid-route. An alternative divergence junction exists upstream of Cart B, where it can reroute *before* committing to the congested path. We compare **Mesh ON** vs. **Mesh OFF** using identical seeds, measuring anticipation time ΔT_{ant} (how many seconds earlier Cart B reroutes), backtracking distance, and total makespan.




Fig. 9: Fleet-size scalability at a fixed crowd of ten pedestrians. Error bars are ± 1 SD, except (c) which shows the Wilson 95% interval.

TABLE V: Mechanism experiment A: V2V mesh anticipation. A corridor is blocked out of line of sight; with the mesh enabled the follower learns of the obstruction before observing it ($N = 50$ paired trials; the same seed drives both arms, so each row is a within-pair comparison). Continuous outcomes use the Wilcoxon signed-rank test and mission success McNemar’s exact test; rows marked *identical* were equal in every pair. p -values are Holm-adjusted across the rows of this table.

Metric	ON	OFF	Test	p (Holm)
Anticipation lead time (s)	10.70 ± 4.20	-0.10 ± 0.04	Wilcoxon	3.8×10^{-9}
Backtrack distance (m)	1.08 ± 0.68	2.73 ± 0.87	Wilcoxon	3.2×10^{-8}
Path length (m)	23.97 ± 3.58	25.81 ± 3.44	Wilcoxon	0.00283
Makespan (s)	21.44 ± 3.40	23.78 ± 3.61	Wilcoxon	0.00196
Mission success	100.0%	100.0%	McNemar	1

The mechanism behaves as designed. With the mesh enabled the trailing cart diverts 10.70 ± 4.20 s before it could have observed the obstruction itself, against -0.10 ± 0.04 s with the mesh disabled—that is, without telemetry the diversion occurs only *after* contact with the blockage. Backtracking falls from 2.73 ± 0.87 m to 1.08 ± 0.68 m and makespan from 23.78 ± 3.61 s to 21.44 ± 3.40 s (Wilcoxon signed-rank, Holm-adjusted $p = 3.8 \times 10^{-9}$ and 2.0×10^{-3} respectively). Success is 100% in both arms: the mesh does not determine *whether* the mission completes here, it determines how much wasted travel is incurred, which is the claim the mechanism was designed to support. Together with the ablation result of Section VI-C, the honest summary is that anticipatory telemetry pays when the topology affords an early alternative and costs a little when it does not.

2) *Experiment B: Corridor Mutex Lock* ($N = 50$ trials): Two carts approach the same single-file corridor from opposite ends, with randomised arrival offsets. The corridor is verified to be genuinely single-file before each trial runs. We compare **Lock ON** against **Lock OFF** on identical seeds.

The effect is large. Mission success rises from 36.0% to 88.0% (McNemar exact, Holm-adjusted $p = 1.5 \times 10^{-4}$) and corridor occupancy falls from 89.41 ± 42.42 s to 40.01 ± 30.87 s ($p = 4.0 \times 10^{-3}$). Without the protocol, the majority of trials fail.

TABLE VI: Mechanism experiment B: distributed corridor mutex. Two carts are routed into the same single-file corridor from opposite ends ($N = 50$ paired trials; the same seed drives both arms, so each row is a within-pair comparison). Continuous outcomes use the Wilcoxon signed-rank test and mission success McNemar’s exact test; rows marked *identical* were equal in every pair. p -values are Holm-adjusted across the rows of this table.

Metric	ON	OFF	Test	p (Holm)
Head-on encounters	1.08 ± 1.18	1.00 ± 0.00	Wilcoxon	1
Deadlocks	0.00 ± 0.00	0.00 ± 0.00	identical	1
Lock wait (s)	0.00 ± 0.00	0.00 ± 0.00	identical	1
Corridor occupancy (s)	40.01 ± 30.87	89.41 ± 42.42	Wilcoxon	0.00403
Makespan (s)	40.05 ± 30.85	89.42 ± 42.39	Wilcoxon	0.00403
Mission success	88.0%	36.0%	McNemar	1.5×10^{-4}

The mechanism is not the one previously claimed, and the distinction matters. An earlier revision described the lock as preventing head-on encounters and eliminating deadlocks. The measurements do not support that description. Head-on encounters are statistically indistinguishable between arms (1.08 ± 1.18 with the lock, 1.00 ± 0.00 without; $p = 1$), and the deadlock counter registers zero in *both* arms—so the $N_{\text{deadlock}} = 0.00$ figure cannot be credited to the lock, since it is equally true without it. Most tellingly, accumulated lock wait time is 0.00 ± 0.00 s: agents essentially never queue at a reserved corridor.

The dataset itself resolves the apparent contradiction. The protocol operates by **cost-projected diversion** rather than by queueing, and the zero wait time is the direct evidence: an agent that deferred by *waiting* would accumulate wait time, and none is recorded in any trial. Instead, when an agent observes a peer’s stronger claim, that claim is projected onto its graph as an effectively infinite edge cost, so D* Lite reroutes through a parallel aisle. The corridor is therefore vacated by re-routing, not by turn-taking—which is consistent with corridor occupancy falling by more than half while wait time stays at zero.

The head-on encounter still occurs, because it is what triggers the exchange of claims; what the protocol changes is the outcome, not the frequency. Correspondingly, the failure mode

without the lock is not a deadlock in the strict sense but a persistent mutual obstruction that expires against the time budget, which is why a counter keyed on unroutable planner states never fires in either arm. Diagnostic instrumentation of individual trials (not part of the released datasets, and therefore not quoted numerically here) confirms the routing divergence directly: with the lock enabled the agents traverse nodes outside the contested corridor and issue roughly twice as many replans, while with it disabled they remain on the contested corridor throughout.

We report this because it changes what the protocol should be credited with. R_{lock} is not a queueing discipline that serialises access to a bottleneck; it is a distributed mechanism for converting a peer's intent into a routing cost, which is a stronger and more general claim, but a different one. It also explains why the corresponding wait-time column is uniformly zero throughout this paper, including in the fleet-scaling results of Section VI-F: there is no queue to measure.

VII. CONCLUSION & FUTURE RESEARCH DIRECTIONS

This paper introduced the Distributed Dynamic Route Optimization (D²RO) framework powered by Socially-Weighted Distributed Graph Optimization (SW-DGO), formalizing a 5-component edge traversal cost function $C(u, v, t) = w_D D + w_M W_{\text{mesh}} + w_H H_{\text{prox}} + w_R R_{\text{lock}} + w_S S_{\text{trolley}}$ solved by incremental D* Lite repair. We assess the outcome against the three aims set out in Section I-D.

A1 (social compliance) is met, as an explicit trade. Across $N = 100$ seed-paired trials D²RO holds median intimate-space exposure at 0 ticks [IQR 0, 0], against 128 for Static A^* and 204 for artificial potential fields ($p < 10^{-15}$). It pays $2.6\times$ the makespan of the socially blind shortest path (47.18 ± 13.40 s vs. 18.00 ± 0.00 s) and matches, rather than exceeds, the best baseline success rate (99.0% vs. 100.0%, a difference of one trial). We regard trading travel time for the near-elimination of personal-space intrusion as the correct engineering choice in human-shared space, and we present it as a trade rather than as dominance.

A2 (distributed coordination) is met, with a corrected account of how. Under controlled conditions the mesh advances rerouting by 10.70 ± 4.20 s and halves backtracking, while the corridor protocol raises success from 36.0% to 88.0%. Both mechanisms are, however, invisible or mildly counterproductive in the broad benchmark, where the topologies they address seldom arise; and the corridor protocol works by projecting a peer's claim into the cost field so that D* Lite diverts, not by queueing agents at the bottleneck. Its contribution is distributed intent propagation, not mutual exclusion in the conventional sense.

A3 (generality and real-time feasibility) is met, within a stated fleet bound. The unchanged planner achieves 99.0%, 100.0% and 95.0% success across supermarket, hospital and airport topologies, with vertex repair in 0.09–0.19 ms: under 0.4% of the 50 ms control budget across every density and fleet size evaluated. Generality does not extend indefinitely with fleet size—success falls to 78.0% at twelve carts in a 36×24 m floorplan as mesh traffic grows super-linearly—and

we state that ceiling as a limitation of the present design rather than describing it as graceful degradation.

A. Theoretical Guarantees & Their Empirical Limits

1) Heuristic Admissibility & Incremental Optimality:

Because the Euclidean distance heuristic $h(s, s_{\text{goal}}) \leq c^*(s, s_{\text{goal}})$ is strictly admissible and consistent, D* Lite guarantees optimal path extraction with respect to the currently observed edge cost field $C(u, v, t)$ while recomputing only the subgraph perturbed by dynamic obstacles [8]. We verified this property empirically against fresh Dijkstra solutions over 150 randomised scenarios comprising 750 successive edge-cost repairs.

2) Conflict Resolution in Single-File Bottlenecks:

Directional reservations broadcast as LOCK_REQUEST packets inflate the contested edge cost for the deferring agent, which then diverts through a parallel aisle or holds in a Turnout Alcove (V_{alcove}). We claim empirical conflict resolution, not deadlock freedom: no formal guarantee is established here, and the zero deadlock counts we measure occur in both the lock-enabled and lock-disabled arms, so they cannot be attributed to the protocol. A liveness proof under bounded latency and message loss remains open.

B. Limitations

We state the principal limitations directly. (i) The weights w_D, w_M, w_H, w_R, w_S were set by hand; no sensitivity analysis over them is reported here, so the reported operating point is not shown to be robust to their perturbation. (ii) The ORCA and Local MAPF baselines are our own implementations and returned 0% success; until they are validated against reference implementations we treat those figures as properties of our code, and no claim in this paper depends on them. (iii) All results are simulation-based; no physical fleet was deployed. (iv) The turn penalty α_{turn} of Eq. (3) is not implemented in the evaluated system, as it requires a heading-augmented search space; the cornering deceleration present in the motion layer is a partial substitute. (v) Fleet scaling is bounded as described above.

C. Considerations for a Physical Deployment

The evaluation in this paper is entirely simulation-based. The following are therefore *requirements a physical deployment would have to satisfy*, not properties of the implemented system; none of the hardware described below forms part of the evaluated framework.

1) RF Signal Attenuation & Steel Fixture Multipath:

Metallic retail shelf gondolas and merchandise packaging attenuate 2.4 GHz RF signals by -15 to -25 dBm. A deployment would likely require Sub-GHz (868/915 MHz) or dedicated short-range (5.9 GHz IEEE 802.11p / DSRC) links with multi-hop TTL forwarding. The exponential penalty decay $W_{\text{mesh}}(t) = W_0 \exp(-\lambda t)$ that is implemented and evaluated here provides the relevant safeguard in principle: stale telemetry expires on

its own, so intermittent packet loss cannot permanently poison the routing graph. Our mesh model includes latency and loss, but has not been validated against measurements from a real RF environment.

- 2) **Indoor Positioning & Sensor Fusion:** All results assume the pose is known. A physical cart would need to recover it—plausibly by fusing wheel encoders, an IMU, Ultra-Wideband trilateration and LiDAR scan-matching through an Extended Kalman Filter [18], [20]—and the resulting localisation error would propagate into both the proxemic field and the corridor reservations. Quantifying that sensitivity is future work.
- 3) **Payload Invariance & Non-Holonomic Kinodynamics:** Service carts undergo large payload shifts (on the order of 15 kg empty to 65 kg loaded), altering the centre of gravity and rotational inertia. The evaluated system uses fixed kinodynamic limits; a deployment would need to modulate $a_{\max}$ and $\kappa_{\max}$ from load sensing to avoid wheel scrubbing or tip-over. This is not implemented or evaluated here.

D. Future Research Directions

The limitations above define the immediate agenda; the last two items are longer-range.

- 1) **Weight Sensitivity & Calibration:** The five weights are currently hand-set. A systematic sweep—perturbing each weight over a multiplicative grid while holding the others fixed, and reporting the resulting success/makespan/social-exposure surface—would establish whether the operating point is a broad plateau or a narrow ridge, and would replace hand-tuning with a stated calibration procedure.
- 2) **Baseline Validation:** Reproducing our ORCA implementation’s behaviour against a reference (e.g. RVO2) on a canonical benchmark, so that comparative claims rest on validated baselines.
- 3) **Mesh Traffic Suppression:** Relevance filtering and spatial scoping of broadcasts, to lift the fleet-size ceiling identified in Section VI-F.
- 4) **Heading-Augmented Search:** Implementing α_{turn} over a state space augmented with heading, allowing turn cost to enter the graph search rather than only the motion layer.
- 5) **Hybrid Learning-Guided Search:** Coupling Graph Neural Networks or deep reinforcement learning (e.g., PRIMAL / Learn to Follow) for global sub-goal allocation with the deterministic D²RO engine for kinodynamic execution.
- 6) **Heterogeneous Multi-Agent Ecosystems:** Extending SW-DGO to mixed fleets—delivery pods, floor-scrubbing AGVs, human-operated wheelchairs—via vehicle-specific agility weights in the R_{lock} reservation tuple.

DATA AND CODE AVAILABILITY

All simulation source code, experimental harnesses, test suites, raw CSV datasets and vector figures are available at <https://github.com/polla-fattah/D2RO>. Every number, table and

TABLE VII: Provenance of the reported results. Every entry marked *complete* is regenerated from the committed raw data by the analysis pipeline; entries marked otherwise are not reported in this manuscript.

Experiment	Dataset status
Comparative benchmark	complete
Component ablation	complete
Cross-domain generalisation	complete
Crowd-density scalability	complete
Fleet-size scalability	complete
Mechanism A: mesh anticipation	complete
Mechanism B: corridor mutex	complete

figure in this paper is produced from the committed raw data by a single analysis pipeline; none is entered by hand. Table VII records the provenance status of each dataset backing the results reported above.

ACKNOWLEDGMENT

The authors thank the Autonomous Systems and Robotics Laboratory for hardware access and computing support.

REFERENCES

- [1] A. F. M. Azlan, N. A. M. Nasir, N. A. I. Nordin, M. D. I. Shaharuddin, R. Ariffin, N. H. I. M. Lazim, and K. Nabilah, “Autonomous mobile trolley robot: Int-cart,” *Semarak Proceedings of Applied Sciences and Engineering Technology*, vol. 1, no. 1, pp. 1–9, 2024.
- [2] T. Kruse, A. K. Pandey, R. Alami, and A. Kirsch, “Human-aware robot navigation: A survey,” *Robotics and Autonomous Systems*, vol. 61, no. 12, pp. 1726–1743, 2013.
- [3] C. Chen, S. Liu, M. Liu, D. Zeng, and D. Manocha, “Relational graph learning for crowd navigation,” *IEEE Robotics and Automation Letters*, vol. 5, no. 2, pp. 2451–2458, 2020.
- [4] J. Van den Berg, M. Lin, and D. Manocha, “Reciprocal velocity obstacles for real-time multi-agent navigation,” *IEEE Transactions on Robotics*, vol. 24, no. 4, pp. 706–718, 2008.
- [5] R. Stern, N. Sturtevant, A. Felner, S. Koenig, H. Ma, T. Walker, J. Li, D. Atzmon, L. Cohen, T. K. S. Kumar, E. Boyarski, and G. Sharon, “Multi-agent pathfinding: Definitions, variants, and benchmarks,” *Symposium on Combinatorial Search (SoCS)*, vol. 10, no. 1, pp. 151–158, 2019.
- [6] S. Dergachev and K. Yakovlev, “Distributed multi-agent navigation based on reciprocal collision avoidance and local multi-agent path finding,” *Robotics and Autonomous Systems*, vol. 145, p. 103856, 2021.
- [7] J. Gielis, A. Subotic, and A. Prorok, “Co-design of communication and path planning for multi-robot teams: A survey and taxonomy,” *IEEE Transactions on Robotics*, vol. 38, no. 5, pp. 2700–2718, 2022.
- [8] S. Koenig and M. Likhachev, “D* Lite,” *Eighteenth National Conference on Artificial Intelligence (AAAI-02)*, pp. 476–483, 2002.
- [9] H. Ma, J. Li, T. K. S. Kumar, and S. Koenig, “Lifelong multi-agent path finding for online pickup and delivery,” *Autonomous Robots*, vol. 41, no. 7, pp. 1613–1630, 2017.
- [10] S. Dergachev and K. Yakovlev, “Decentralized target and priority swapping for unlabeled multi-agent path finding,” *IEEE Robotics and Automation Letters*, vol. 9, no. 3, pp. 2451–2458, 2024.
- [11] M. Keskin, M. Guler, and S. Sen, “Automated negotiation-based decentralized multi-agent path finding in shared environments,” *IEEE Transactions on Intelligent Vehicles*, vol. 9, no. 2, pp. 3102–3114, 2024.
- [12] G. Sartoretti, J. Kerr, Y. Shi, G. Wagner, T. K. S. Kumar, S. Koenig, and H. Choset, “PRIMAL: Pathfinding via reinforcement and imitation multi-agent learning,” *IEEE Robotics and Automation Letters*, vol. 4, no. 3, pp. 2378–2385, 2019.
- [13] A. Skrynnik, A. Andreychuk, A. Panov, and K. Yakovlev, “Learn to follow: Decentralized lifelong multi-agent pathfinding with guidance graphs,” *Autonomous Agents and Multi-Agent Systems*, vol. 38, no. 1, p. 12, 2024.
- [14] K. Al-Mutib, E. Mattar, M. Alsulaiman, and H. Ramdane, “D* Lite based real-time path planning for multi-robot systems,” *International Conference on Advanced Robotics (ICAR)*, pp. 145–151, 2012.

- [15] G. Wagner and H. Choset, “M*: Subdimensional expansion for multirobot path planning,” *IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, pp. 3953–3958, 2011.
- [16] V. Slyusar and M. Kulich, “Evaluation of routing protocols for mobile ad-hoc networks in multi-robot exploration,” *Journal of Intelligent & Robotic Systems*, vol. 84, no. 1, pp. 215–231, 2016.
- [17] P. Edwige, P.-Y. Lajoie, and G. Beltrame, “Swarm-SLAM: Sparse decentralized collaborative simultaneous localization and mapping with ad-hoc wireless communication,” *IEEE Transactions on Robotics*, vol. 40, pp. 1120–1136, 2024.
- [18] F. Zafari, A. Gkelias, and K. K. Leung, “A survey of indoor localization systems and technologies,” *IEEE Communications Surveys & Tutorials*, vol. 21, no. 3, pp. 2568–2599, 2019.
- [19] C. Clark, S. Rock, and B. Williams, “TEAM: Topological mapping and localization using ad-hoc localized multi-robot networks,” *Autonomous Robots*, vol. 45, no. 4, pp. 521–538, 2021.
- [20] A. Nugraha, B. Pratama, and T. Wibowo, “Sensor fusion of ultra-wideband and wheel odometry using extended kalman filters for indoor autonomous navigation,” *IEEE Sensors Journal*, vol. 24, no. 4, pp. 4512–4521, 2024.
- [21] X. Yao, Y. Tan, D. Gu, and H. Tang, “HA-VLN: A benchmark for human-aware vision-and-language navigation with path diversity and generalization,” in *Proceedings of the 32nd ACM International Conference on Multimedia (ACM MM)*. ACM, 2024, pp. 1–10.